

Compiling and running code on CPU nodes

Xin Li

Cray programming environment (CPE)

Reference page: [Compilers and libraries](#)

The Cray Programming Environment (CPE) provides consistent interface to multiple compilers and libraries.

- In practice, we recommend
 - `ml cpeCray/23.03`
 - `ml cpeGNU/23.03`
 - `ml cpeA0CC/23.03`
- The `cpeCray` , `cpeGNU` and `cpeA0CC` modules are available after `ml PDC/23.03`
- No need to `module swap` or `module unload`

Compiler wrappers

- Compiler wrappers for different programming languages
 - `cc` : C compiler wrapper
 - `CC` : C++ compiler wrapper
 - `ftn` : Fortran compiler wrapper
- The wrappers choose the required compiler version and target architecture options.
- Automatically link to MPI library and math libraries
 - MPI library: `cray-mpich`
 - Math libraries: `cray-libsci` and `cray-fftw`

Compile a simple MPI code

- `hello_world_mpi.f90`

```
program hello_world_mpi
include "mpif.h"
integer myrank,mysize,ierr
call MPI_Init(ierr)
call MPI_Comm_rank(MPI_COMM_WORLD,myrank,ierr)
call MPI_Comm_size(MPI_COMM_WORLD,mysize,ierr)
write(*,*) "Processor ",myrank," of ",mysize,": Hello World!"
call MPI_Finalize(ierr)
end program
```

```
ftn hello_world_mpi.f90 -o hello_world_mpi.x
```

What flags does the `ftn` wrapper activate?

- Use the flag `-craype-verbose`

```
ftn -craype-verbose hello_world_mpi.f90 -o hello_world_mpi.x
```

Compile a simple MPI code

```
user@uan01:> srun -n 8 ./hello_world_mpi.x
Processor      4 of      8 : Hello World!
Processor      6 of      8 : Hello World!
Processor      7 of      8 : Hello World!
Processor      0 of      8 : Hello World!
Processor      1 of      8 : Hello World!
Processor      2 of      8 : Hello World!
Processor      3 of      8 : Hello World!
Processor      5 of      8 : Hello World!
```

Compile a simple linear algebra code

[Link to the code](#)

Use cray-libsci

```
ml PDC/23.03 cpeGNU/23.03
```

```
cc dgemm_test.c -o dgemm_test_craylibsci.x
```

Compile a simple linear algebra code

Use openblas

```
ml openblas/0.3.24-gcc-s34
```

```
export ROOTOPENBLAS=/pdc/software/23.03/spack/openblas-0.3.24-s34z4q2  
cc dgemm_test.c -o dgemm_test_openblas.x -I$ROOTOPENBLAS/include -L$ROOTOPENBLAS/lib -lopenblas
```


Check the linked libraries

```
ldd dgemm_test_craylibsci.x
```

```
ldd dgemm_test_openblas.x
```

Check the linked libraries

```
ldd dgemm_test_craylibsci.x
```

```
...  
libsci_gnu_82.so.5 => /opt/cray/pe/lib64/libsci_gnu_82.so.5  
...
```

```
ldd dgemm_test_openblas.x
```

```
...  
libopenblas.so.0 => /.../openblas-0.3.24-s34z4q2/lib/libopenblas.so.0  
...
```

Exercise: Compile and run the dgemm_test code

- Run on a single core in the `shared` partition

```
salloc -n 1 -t 10 -p shared -A <name-of-allocation>  
srun -n 1 ./dgemm_test_craylibsci.x  
srun -n 1 ./dgemm_test_openblas.x  
exit
```

- Expected output:

2.700	4.500	6.300	8.100	9.900	11.700	13.500
4.500	8.100	11.700	15.300	18.900	22.500	26.100
6.300	11.700	17.100	22.500	27.900	33.300	38.700

Exercise: Compile and run `fftw_test` code

```
ml cray-fftw/3.3.10.3

wget https://people.math.sc.edu/Burkardt/c_src/fftw/fftw_test.c

cc --version
cc fftw_test.c -o fftw_test.x

ldd fftw_test.x

salloc -n 1 -t 10 -p shared -A <name-of-allocation>
srun -n 1 ./fftw_test.x
```

Compilation of large program

- Examples at <https://www.pdc.kth.se/software>

Environment variables for manual installation of software

- Environment variables for compilers

```
export CC=cc  
export CXX=CC  
export FC=ftn  
export F77=ftn
```

- Environment variables for compiler flags
 - add `-I` , `-L` , `-l` , etc. to Makefile
- Environment variables at runtime
 - prepend to `PATH` , `LD_LIBRARY_PATH` , etc.

What happens when loading a module

```
ml show openblas/0.3.24-gcc-s34
```

```
whatis("OpenBLAS: An optimized BLAS library")
prepend_path("PATH", "/pdc/software/23.03/spack/openblas-0.3.24-s34z4q2/bin")
prepend_path("LD_LIBRARY_PATH", "/pdc/software/23.03/spack/openblas-0.3.24-s34z4q2/lib")
prepend_path("PKG_CONFIG_PATH", "/pdc/software/23.03/spack/openblas-0.3.24-s34z4q2/lib/pkgconfig")
prepend_path("CMAKE_PREFIX_PATH", "/pdc/software/23.03/spack/openblas-0.3.24-s34z4q2/.")
...
```

When running your own code

- Load correct programming environment (e.g. `cpeGNU`)
- Load correct dependencies (e.g. `openblas` if your code depends on it)
- Properly prepend to environment variables (e.g. `PATH` , `LD_LIBRARY_PATH`)
- Choose correct SLURM settings

SLURM settings for hybrid MPI/OpenMP code

- `--nodes` number of nodes
- `--ntasks-per-node` number of MPI processes
- `--cpus-per-task` 2 x number of OpenMP threads (because of SMT)
- `OMP_NUM_THREADS` number of OpenMP threads
- `OMP_PLACES` cores

Example job script

- 64 MPI x 2 OMP per node (main partition)

```
#!/bin/bash

#SBATCH -A ...
#SBATCH -J my_job
#SBATCH -t 01:00:00
#SBATCH -p main

#SBATCH --nodes=2
#SBATCH --ntasks-per-node=64
#SBATCH --cpus-per-task=4

module load ...

export OMP_NUM_THREADS=2
export OMP_PLACES=cores

srun ...
```

Example job script

- 2 MPI x 2 OMP per node (shared partition)

```
#!/bin/bash

#SBATCH -A ...
#SBATCH -J my_job
#SBATCH -t 01:00:00
#SBATCH -p shared

#SBATCH --ntasks=2
#SBATCH --cpus-per-task=4

module load ...

export OMP_NUM_THREADS=2
export OMP_PLACES=cores

srun ...
```

Exercise: Hybrid MPI/OpenMP code for matrix-matrix multiplication

- Preparation

```
mkdir -p matmul_test && cd matmul_test
```

- Copy python code [matmul_mpi_omp_test.py](#) to the same folder

Exercise: Hybrid MPI/OpenMP code for matrix-matrix multiplication

- Copy job script [job-n1.sh](#)
 - for running on 1 MPI processes with different number of OpenMP threads
- Copy job script [job-n2.sh](#)
 - for running on 2 MPI processes with different number of OpenMP threads
- Submit two jobs

Exercise: Hybrid MPI/OpenMP code for matrix-matrix multiplication

- Result

Setting	Timing
1 MPI x 16 OMP	Time spent in matmul: 2.307 sec
1 MPI x 8 OMP	Time spent in matmul: 3.924 sec
1 MPI x 4 OMP	Time spent in matmul: 6.626 sec
2 MPI x 8 OMP	Time spent in matmul: 2.034 sec
2 MPI x 4 OMP	Time spent in matmul: 3.287 sec
2 MPI x 2 OMP	Time spent in matmul: 6.188 sec