



# Accelerator Programming with OpenMP

# Introduction

- In a perfect world we would like the OpenMP work-sharing directives to automatically run on an accelerator

Host

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char **argv) {

    const int n = 100;
    double *a = (double *) malloc(n * sizeof(double));
    double *b = (double *) malloc(n * sizeof(double));
    double *c = (double *) malloc(n * sizeof(double));

    for (int i = 0; i < n; i++) {
        a[i] = 1 + i;
        b[i] = 1 - i;
    }
    #pragma omp parallel for
    for (int i = 0; i < n; i++) {
        c[i] = a[i] + b[i];
    }

    for (int i = 0; i < 5; i++) {
        printf("c[%d] = %g\n", i, c[i]);
    }
    printf("c[%d] = %g\n", n-1, c[n-1]);

    free(a);
    free(b);
    free(c);
}
```

Thread1

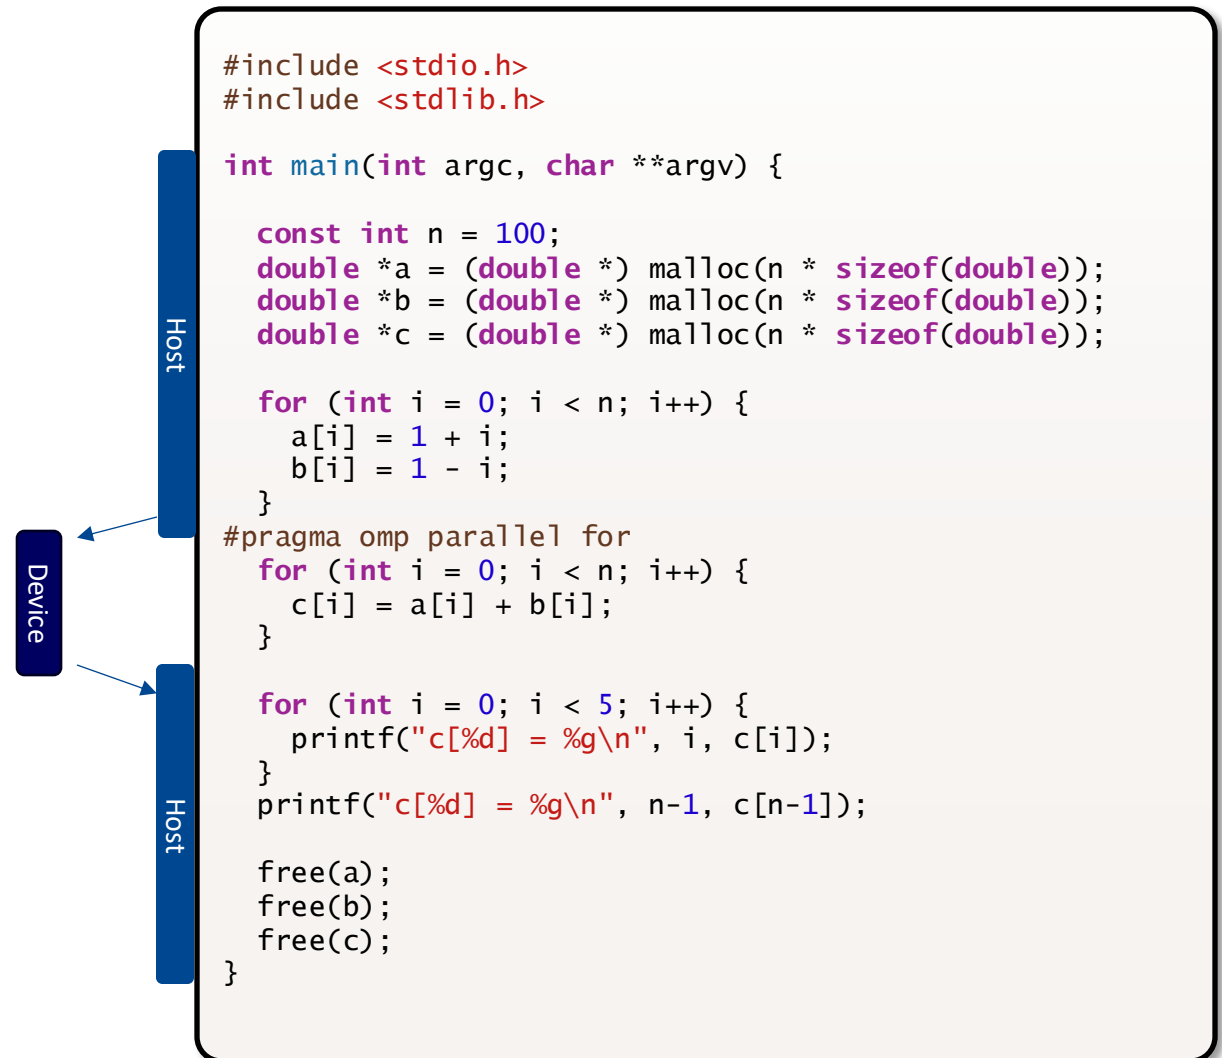
Thread2

Thread3

Thread4



# Introduction

- In a perfect world we would like the OpenMP work-sharing directives to automatically run on an accelerator



# Introduction

- In a perfect world we would like the OpenMP work-sharing directives to automatically run on an accelerator
- OpenMP offload constructs are a set of directives for C/C++ and Fortran to offload work to an accelerator
- The **target** construct offloads the enclosed code to the accelerator and **map** transfers data
- But how is it executed?



```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char **argv) {

    const int n = 100;
    double *a = (double *) malloc(n * sizeof(double));
    double *b = (double *) malloc(n * sizeof(double));
    double *c = (double *) malloc(n * sizeof(double));

    for (int i = 0; i < n; i++) {
        a[i] = 1 + i;
        b[i] = 1 - i;
    }



    #pragma omp target map(a[0:n],b[0:n],c[0:n])
    for (int i = 0; i < n; i++) {
        c[i] = a[i] + b[i];
    }

    for (int i = 0; i < 5; i++) {
        printf("c[%d] = %g\n", i, c[i]);
    }
    printf("c[%d] = %g\n", n-1, c[n-1]);

    free(a);
    free(b);
    free(c);
}
```

# Introduction

- In a perfect world we would like the OpenMP work-sharing directives to automatically run on an accelerator
- OpenMP offload constructs are a set of directives for C/C++ and Fortran to offload work to an accelerator
- The **target** construct offloads the enclosed code to the accelerator and **map** transfers data
- But how is it executed?
- Without additional directives expressing parallelism the kernel will run in **serial**!



```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char **argv) {

    const int n = 100;
    double *a = (double *) malloc(n * sizeof(double));
    double *b = (double *) malloc(n * sizeof(double));
    double *c = (double *) malloc(n * sizeof(double));

    for (int i = 0; i < n; i++) {
        a[i] = 1 + i;
        b[i] = 1 - i;
    }

    #pragma omp target map(a[0:n],b[0:n],c[0:n])
    for (int i = 0; i < n; i++) {
        c[i] = a[i] + b[i];
    }


    for (int i = 0; i < 5; i++) {
        printf("c[%d] = %g\n", i, c[i]);
    }
    printf("c[%d] = %g\n", n-1, c[n-1]);

    free(a);
    free(b);
    free(c);
}
```


# OpenMP offload target parallelism

Several different combinations of device execution directives (we will only cover some of them)


`omp target`




`omp target teams`



`omp target teams  
distribute`



`omp target teams  
distribute parallel for/do`



# OpenMP offload target parallelism

- The **target** construct offloads the code to the accelerator
- The **teams** construct creates a set of teams (but still no work-sharing within the teams)
- With the **distribute** construct work can be distributed between teams (e.g. **target teams distribute**), but still no work-sharing within the teams
- Using a **parallel for/do** construct, work can be distributed within each team (e.g. **target teams distribute parallel for/do**)
- Or use the modern (OpenMP 5) construct **omp target teams loop** to use the full GPU (similar to the **target teams distribute parallel for/do** construct)

# OpenMP device memory

Controlling and minimising data movement to/from the device is the most important step when optimising an OpenMP offloading program

- The **map** clause can be used to control data movement
- Various form of the clause:
  - `map(to:list)` copy data to the device when entering the region
  - `map(from:list)` copy data back to the host when leaving the region
  - `map(tofrom:list)` combination of the above to and from clause
  - `map(alloc:list)` allocate data when entering the region

**Note:** In Fortran nothing special needs to be done to move dynamically allocated memory, while in C/C++ the dimension has to be specified e.g. `map(to:a[0:n])`

**Hint:** Vendors often provide tools to show data movement, on a HPE Cray set the environment variable `CRAY_ACC_DEBUG=1, 2 or 3`



# OpenMP device memory

The **target data** construct creates a data region, to keep the data on the device between different target offload regions (use **map** to control data movement)

```
!$omp target data map(a,b) <- move data to device

!$omp target
  do a lot of work on the device with a,b

do something on the host

!$omp target
  do a lot of work on the device with a,b

!$omp end target data <- move data back to host
```

Update data to/from a device using the **target update to(list)/from(list)** clause

# References

OpenMP specifications:

<http://www.openmp.org/specifications/>

OpenMP API syntax reference guide

<https://www.openmp.org/wp-content/uploads/OpenMPRefGuide-5.2-Web-2024.pdf>